Politecnico di Torino

Master of Science (M.Sc.) in Computer Engineering

Academic Year 2025/2026

Graduation Session March 2026

Design and Implementation of a Cloud-Native Reactive Architecture

The Smart Travel Platform Case Study

Supervisor

Prof. Giovanni Malnati

Candidate

Federico Palazzi

Abstract

The rapid evolution of enterprise web applications has driven a paradigm shift from monolithic architectures to cloud-native *microservices* solutions. This transition is motivated by the need for greater scalability, resilience, and agility in the software development life cycle. This thesis presents the design, development and deployment of "*Smart Travel*", a comprehensive e-commerce platform dedicated to the sale of travel packages. The project serves as a proof of concept (PoC) developed within an industrial internship context, aiming to analyze the complexities and benefits of a modern and distributed software architecture.

The primary objective of this work is not merely the implementation of functional requirements (such as catalog management and order processing), but rather the critical evaluation of competing technologies and the automation of the release process. A central element of this research is the comparative analysis between two leading Java ecosystems: the established *Spring Boot* framework and the cloud-native *Quarkus* framework. Specifically, the project adopts a fully reactive approach for both of them, utilizing *Spring WebFlux* and *Quarkus Reactive* capabilities.

From an architectural perspective, the Smart Travel system is designed to ensure loose coupling and high availability. The backend adopts a microservices pattern supported by *RabbitMQ*, a message broker that facilitates asynchronous *event-driven* communication among services. Data persistence is managed through *MongoDB*, a NoSQL database chosen for its schema flexibility, which is particularly advantageous for handling diverse and evolving travel package structures compared to rigid relational models. On the client side, the user interface is delivered via a *Single Page Application (SPA)* developed with *Angular* and leveraging *RxJS* for reactive state management, ensuring a responsive and fluid user experience.

A significant portion of this thesis is dedicated to the "industrialization" of the software, bridging the gap between development and operations (*DevOps*). The project moves beyond local execution by implementing a robust *Continuous Integration and Continuous Delivery (CI/CD)* pipeline using *Jenkins*. This pipeline automates the entire delivery chain, including code checkout, compilation, container image creation and pushing to a private registry. The final deployment target is *OKD* (the community distribution of *Kubernetes* that powers *Red Hat OpenShift*). The thesis details the orchestration of containerized applications within the OKD cluster, managing critical aspects such as service discovery, route exposure and configuration management via `ConfigMaps` and `Secrets`.

In conclusion, the Smart Travel project demonstrates the viability and robustness of a heterogeneous, fully reactive microservices architecture. It successfully integrates different modern technologies into a cohesive ecosystem, proving that the adoption

of cloud-native principles, while introducing complexity in orchestration and consistency management, offers superior long-term benefits in terms of maintainability and scalability. The work provides a practical blueprint for modernizing legacy Java applications, offering concrete evidence of how the synergy between Reactive frameworks, Event-Driven patterns, and automated DevOps practices defines the state of the art in contemporary software engineering.

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Tristique senectus et netus et malesuada fames ac turpis. Pretium nibh ipsum consequat nisl vel pretium lectus quam. Urna molestie at elementum eu facilisis sed odio morbi quis. Sed egestas egestas fringilla phasellus faucibus scelerisque eleifend. At in tellus integer feugiat. Mauris rhoncus aenean vel elit scelerisque mauris pellentesque pulvinar. Commodo sed egestas egestas fringilla. Nunc lobortis mattis aliquam faucibus purus in massa. Facilisis magna etiam tempor orci eu lobortis elementum nibh. Elementum curabitur vitae nunc sed velit dignissim. Neque volutpat ac tincidunt vitae. Massa id neque aliquam vestibulum morbi blandit cursus risus at. Porta non pulvinar neque laoreet suspendisse interdum consectetur. Turpis in eu mi bibendum. Ut tristique et egestas quis ipsum suspendisse. Integer quis auctor elit sed vulputate mi sit amet. Viverra nam libero justo laoreet sit amet cursus.

Table of Contents

List of Figures VI

Glossary VIII

1 Introduction 1
- 1.1 Context . . . 1
- 1.2 Goal . . . 2
  - 1.2.1 Project description . . . 3
  - 1.2.2 Expected Results . . . 3
- 1.3 Thesis structure . . . 4

2 Background and Technologies 6

3 System Architecture 7

4 Application Design and Implementation 8

5 Infrastructure Automation and Cloud Deployment 9

6 Conclusion 10

Bibliography 11

List of Figures

1.1 Monolithic vs. Microservices architecture 2

Glossary

CI/CD
: Continuous Integration / Continuous Delivery

DX
: Developer Experience

NoSQL
: Not Only SQL

PoC
: Proof of Concept

RxJS
: Reactive Extensions for JavaScript

SPA
: Single Page Application

Chapter 1

Introduction

1.1 Context

The rapid evolution of enterprise web applications over the last decade has driven a significant paradigm shift in software engineering, moving away from traditional monolithic architectures towards cloud-native microservices solutions.

Historically, monolithic applications were the industry standard and they integrated user interface, business logic, and data access layers into a single deployable unit. They offered simplicity during the early stages of development, straightforward local testing, and centralized debugging. However, as applications grow in size, complexity, and user base, the inherent limitations of the monolithic approach become increasingly apparent [1].

A monolithic design relies on tightly coupled components. Consequently, a minor update, a new feature addition, or a simple bug fix requires the entire massive codebase to be recompiled, tested, and redeployed. This rigid structure often leads to development bottlenecks and significantly increases the risk of system-wide failures during deployments. Furthermore, scaling a monolith is inherently inefficient; if a single specific component (such as the payment processing module) experiences a surge in traffic, the entire monolithic application must be scaled horizontally, leading to unnecessary resource consumption and inflated operational costs.

To overcome these scalability and maintainability challenges, the industry has widely adopted the microservices architecture [1]. This modern paradigm decomposes a large application into a suite of small, autonomous, and loosely coupled services. Each microservice is organized around a specific business capability, encapsulates its own data, and communicates with other services through lightweight, well-defined application programming interfaces (APIs) or asynchronous message brokers. This decentralized approach offers superior modularity, improved fault isolation - the

failure of one service does not inevitably crash the entire system - and the agility to deploy and scale individual services independently based on their specific demand. However, the transition to microservices is not without its own set of complexities. It introduces challenges such as increased operational overhead, the need for sophisticated monitoring and logging solutions, and the complexity of managing distributed transactions and data consistency across services. Despite these challenges, the microservices architecture has become the de facto standard for building scalable, resilient, and maintainable enterprise applications in the cloud era.

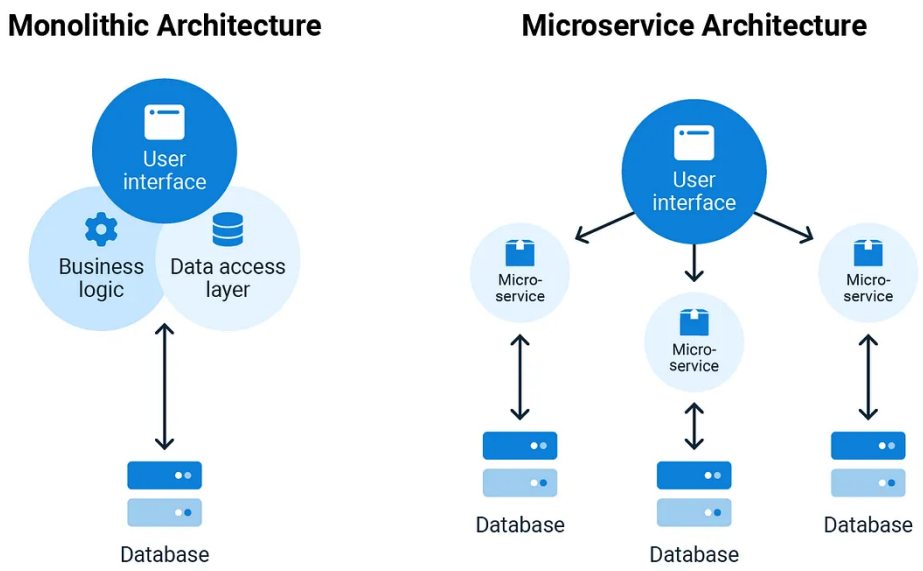

Figure 1.1: Monolithic vs. Microservices architecture. *Source: [2]*

1.2 Goal

The primary objective of this thesis is to explore, apply, and critically evaluate modern cloud-native technologies and microservices architectures. To achieve this, the "*Smart Travel*" platform was designed and developed as a proof of concept (PoC). Rather than simply focusing on the implementation of functional business requirements, this project serves primarily as a practical testbed to analyze the complexities, challenges, and benefits of building a modern, distributed software system. The project was conducted during an internship at a software company, situated directly within the industry's broader context of architectural modernization.

1.2.1 Project description

The Smart Travel application was designed as a comprehensive e-commerce platform dedicated to the sale of travel packages. The system manages the entire customer journey, from the public consultation of a catalog containing flights, accommodations, and organized activities, to the execution of the purchasing process, order history management, and the automated dispatch of email notifications.
A central technical goal of this research is the critical evaluation and comparative analysis of two leading Java ecosystems used to build the platform's microservices: the established *Spring Boot* framework and the newer cloud-native *Quarkus* framework. Specifically, the project adopts a fully reactive programming paradigm for both frameworks, utilizing *Spring WebFlux* and *Quarkus Reactive* capabilities. This enables non-blocking, asynchronous execution, allowing the system to handle a massive number of concurrent requests with a highly optimized `thread pool`.
Furthermore, this thesis aims to explore the "industrialization" of software through the adoption of *DevOps* practices. The objective is to move beyond local execution by implementing a robust *Continuous Integration and Continuous Delivery* (CI/CD) pipeline. By automating the entire delivery chain - from code checkout and compilation to container image creation and pushing to a private registry - the project demonstrates how to bridge the gap between development and operations.
Finally, the thesis details the deployment and orchestration of these containerized microservices within an *OKD* cluster (the community distribution of *Red Hat OpenShift*), managing critical operational aspects such as service discovery, route exposure, and configuration management.

1.2.2 Expected Results

The execution of this thesis work is expected to produce both concrete technical artifacts and valuable analytical insights. The expected results can be categorized into three main areas:

1. From a software engineering perspective, the project aims to deliver a fully functional PoC of the e-commerce platform. This includes the successful implementation of the frontend interface and the backend microservices, validating the practical application of asynchronous communication via RabbitMQ and flexible data management using MongoDB. The resulting system is expected to demonstrate the resilience and scalability inherent in a properly designed cloud-native architecture.

2. From an analytical standpoint, the research provides a qualitative comparative evaluation between Spring Boot and Quarkus within a reactive programming context. Given the scope of the internship, the comparison focuses heavily

on the *Developer Experience* (DX) rather than quantitative runtime metrics. The expected outcome is a clear assessment of the differences in syntax, configuration management, boilerplate code reduction, and the overall ease of implementing reactive paradigms in both frameworks. This comparison will offer valuable insights into the practical development lifecycle, ecosystem maturity, and the learning curve associated with adopting Quarkus compared to the traditional Spring environment.

3. From an operational perspective, the expected result is the successful demonstration of a fully automated DevOps lifecycle. By delivering a working CI/CD pipeline via Jenkins that culminates in a live deployment on an OKD cluster, the project will prove the viability and necessity of infrastructure automation. This includes the ability to manage containerized applications at scale, handle configuration and secrets management effectively, and ensure seamless service discovery and routing in a production-like environment.

1.3 Thesis structure

The thesis is structured as follows:

- **Chapter 2: Background and Technologies** - Provides an overview of the core concepts and technologies underpinning the project, including microservices and event-driven architectures, imperative and reactive programming paradigms, backend and frontend frameworks (Spring Boot, Quarkus, Angular), NoSQL databases (MongoDB), and API communication models (REST vs. GraphQL).

- **Chapter 3: System Architecture** - Details the architectural decisions, technology stack, and design patterns used in the Smart Travel project. It covers the microservices decomposition, the event-driven communication via message brokers, and the frontend design.

- **Chapter 4: Application Design and Implementation** - Offers a detailed description of the development process. This includes data modeling, API contract definitions, and explanations of key software components, accompanied by code snippets comparing the Spring Boot and Quarkus implementations.

- **Chapter 5: Infrastructure Automation and Cloud Deployment** - Describes the different steps of the CI/CD pipeline, Jenkins automation configurations, and the final orchestration of containerized applications within the OKD cluster environment.

- **Chapter 6: Conclusion** - Summarizes the project findings, evaluates the results of the technology comparison, discusses the limitations encountered during development, and outlines future work directions for the platform.

Chapter 2

Background and Technologies

2.1 Microservices Architecture: Principles, benefits, and challenges.
2.2 Programming Paradigms: Imperative vs. Reactive programming models.
2.3 Event-Driven Architecture: Asynchronous communication and message brokers (focus on RabbitMQ).
2.4 Backend Technologies: Spring Boot and Quarkus
2.5 Frontend Technologies: Single Page Applications (SPA), Angular, RxJS, and PrimeNG.
2.6 Database Technologies: NoSQL concepts and MongoDB.
2.7 API Paradigms: REST vs. GraphQL.

Chapter 3

System Architecture

3.1 Requirements Analysis: Functional and non-functional requirements of the Smart Travel platform.
3.2 High-Level Architecture: Decomposing the system into microservices (Users, Catalog, Orders, Notifications).
3.3 Data Modeling: Designing the MongoDB collections and handling flexible schemas.
3.4 Communication Patterns: * 3.4.1 Synchronous communication (REST/GraphQL)
3.4.2 Asynchronous communication (RabbitMQ events)

Chapter 4

Application Design and Implementation

4.1 Backend Development: * 4.1.1 Implementing the reactive services with Spring Boot
4.1.2 Implementing the reactive services with Quarkus
4.2 Frontend Development: Building the user interface, state management with RxJS, and connecting to the backend.
4.3 Core Workflows:
4.3.1 The catalog search and filtering mechanism
4.3.2 The payment integration flow
4.3.3 The event-driven notification system

Chapter 5

Infrastructure Automation and Cloud Deployment

5.1 The DevOps Culture: Bridging the gap between development and operations.
5.2 Continuous Integration & Continuous Delivery (CI/CD): Automating the build, test, and containerization process using Jenkins.
5.3 Container Orchestration with Kubernetes/OKD: * 5.3.1 OKD Cluster overview
5.3.2 Workload management: Pods and Deployments
5.3.3 Configuration and Security: ConfigMaps and Secrets
5.3.4 Networking: Services and Route exposure

Chapter 6

Conclusion

6.1 Technology Evaluation: Spring Boot vs. Quarkus: Comparing developer experience, startup times, resource consumption, and reactive capabilities.
6.2 Architectural Evaluation: Assessing the scalability, resilience, and maintainability of the implemented microservices and event-driven patterns.
6.3 Testing and Validation: Ensuring the functional requirements of the Smart Travel application were met.

Bibliography

[1] Muskaan Goyal and Pranav Bhasin. *Moving from monolithic to microservices architecture for multi-agent systems*. 2025. URL: https://arxiv.org/pdf/2505.07838 (cit. on p. 1).

[2] CoderCo. *Monoliths vs Microservices: A Guide to Choosing the Right Architecture for Your Application*. 2023. URL: https://blog.coderco.io/p/monoliths-vs-microservices-a-guide (cit. on p. 2).